

Finite-Blocklength Bounds with Prior Optimization: the Φ -view

Nir Elkayam
the fbl library

Abstract

fbl is a validated library for one-shot / finite-blocklength information theory built on the pairwise-error-probability (PEP) error-spectrum framework. It covers channel coding, rate-distortion (average and excess), and joint source-channel coding (JSCC), in two flavours (exact one-shot / method-of-types). Its distinguishing feature is the *achievability* prior optimization, which we show is a single *convex program* on the simplex, $\max_Q c^\top \Phi(AQ)$, where the kernel of the bound chooses the fixed potential Φ . We solve it directly by a first-order *simplex march* that is exact for any kernel and certifies its own optimality by a water-filling (KKT) condition; the classical exact convex solvers (a quadratic program for the RCU⁺ kernel, a bracketing LP otherwise) are special cases retained as validation anchors. We report a self-contained figure suite for four pinned use cases, all numbers backed by a 107-test cross-check suite.

1 Introduction

For each setting the library computes three quantities: the *achievable* (random-coding) bound, the *converse* (meta-converse LP), and—one-shot only—a *Monte-Carlo* realization. The converse optimizes the prior point-wise (a linear program at one threshold). The achievability bound is a kernel integral of the *whole* error spectrum, so its prior optimization is global. We show it is nonetheless a convex program with a common form, and that the optimal prior can be found exactly and certified intrinsically. The practical question the suite answers is then sharp: *given the certified optimal prior, how large is the gap to a memoryless prior?*

2 The unified Φ -view

Fix an output y and a metric $m(x, y) \geq 0$. Under the dithered tie rule the PEP of a candidate drawn from the prior Q_X is uniform on $[0, 1]$, so the value-weighted spectrum and its average against a non-negative kernel κ are

$$F(w \mid y) = \mathbb{E}_{X \sim Q_X} [m(X, y) \mathbf{1}\{\text{PEP} \leq w\}], \quad J = \int_0^1 F(w \mid y) \kappa(w) dw. \quad (1)$$

The prior enters through one linear map—the cumulative-rank profile $(AQ_X)(\nu) = Q_X\{x : m(x, y) \geq \nu\}$ —and the value through the functional $c^\top v = \int_0^\infty v(\nu) d\nu$. With the kernel potential $\Phi_\kappa(\sigma) = \int_0^1 \min\{w, \sigma\} \kappa(w) dw$ one has the *nonlinear program of the PEP*,

$$\boxed{J(Q_X) = c^\top \Phi_\kappa(AQ_X)} \quad \text{finite alphabet: } J = \sum_j c_j \Phi_\kappa(\sigma_j), \quad (2)$$

with σ_j the cumulative masses (sorted by metric) and c_j the value-gaps. The potential satisfies $\Phi_\kappa(0) = 0$, $\Phi'_\kappa(\sigma) = \int_\sigma^1 \kappa \geq 0$, $\Phi''_\kappa = -\kappa$, so Φ_κ is concave; since AQ_X is linear and $c \geq 0$ in the matched case, J is **concave in Q_X** and prior optimization is a convex program over the simplex. Two tails occur: the *error* tail (channel/JSCC) maximizes a concave J ; the *correct* tail (rate-distortion) minimizes a convex J in the dual (CDF) variable.

The kernel chooses the program. Table 1 lists the potentials. The converse is the point-mass kernel $\kappa = \delta_{w_0}$, whose ramp potential is piecewise-linear (an LP). The RCU⁺ potential is a clamped parabola—*quadratic*—so its convex program is a quadratic program (QP). The exact-RC and rate-distortion potentials are non-polynomial, handled either by a bracketing LP (secant/tangent of Φ) or exactly by the march below.

bound	potential Φ_κ	shape	program
converse, level w_0	$\min\{\sigma, w_0\}$	piecewise-linear	LP
channel, RCU ⁺	$\sigma - \frac{M-1}{2}\sigma^2$ (clamped)	quadratic	QP
channel, exact	$\frac{1}{M}(1 - (1 - \sigma)^M)$	concave	bracket / march
rate-distortion, exact	$(1 - \tau)^M$	convex	bracket / march
rate-distortion, smooth	$e^{-M\tau}$	convex	bracket / march

Table 1: The potentials instantiating (2) (M the codebook size). The QP is simply the case of a *quadratic* Φ .

The simplex march and its certificate. We maximize $f = \text{sense} \cdot J$ (sense = +1 channel/JSCC, −1 RD) directly on the simplex by projected gradient / Frank–Wolfe, using the analytic water-fill gradient

$$\nabla f(Q) = A^\top(c \odot \Phi'(AQ)), \quad g(x) = \sum_y \sum_{i \text{ fed by } x} \text{ratio}_i \sum_{j \geq i} c_j \Phi'(\sigma_j). \quad (3)$$

The prior obeys a *product* of simplices (one global simplex for channel/RD; one per source-type block for JSCC). Optimality is the **water-filling KKT condition**— ∇f flat on the support and dominated off it, per block—and the Frank–Wolfe gap $g(Q) = \sum_b (\max_b \nabla f - \langle \nabla f_b, Q_b \rangle)$ is a rigorous suboptimality certificate, $0 \leq f(Q^*) - f(Q) \leq g(Q)$. Thus the achieved bound is always valid (a feasible prior), and $g(Q)$ brackets the prior-optimized optimum without any second solver. The march is exact for every kernel; the QP/bracketing-LP are cvxpy instantiations of the *same* program, kept as independent machine-precision anchors.

3 The library

The code is a small dependency DAG (infra $\rightarrow$ engines $\rightarrow$ prior-opt $\rightarrow$ examples). The prior-opt layer is `phi_view` (the relaxation $J = c^\top \Phi(AQ)$, its potentials and evaluators) and `phi_simplex` (the march and the `check_kkt` certificate); the exact convex solvers remain as anchors.




Figure 1: Module dependency DAG. Prior optimization is the Φ -view mechanism (`phi_view`/`phi_simplex`) with the exact solvers as anchors.




Figure 2: The unified program: the kernel selects Φ (Dirac $\rightarrow$ LP, $\text{RCU}^+ \rightarrow$ QP, general $\rightarrow$ bracket), and the simplex march solves all of them exactly.

Validation. No closed-form oracle exists, so every quantity is cross-checked (107 tests): one-shot $\leftrightarrow$ type-based at small n ; RCU expectation $\leftrightarrow$ Monte-Carlo; converse $\leq$ achievable; the Φ -view identity $J = c^\top \Phi(AQ)$ against an independent computation in both representations and all settings; and the march optimum against the exact solvers *and* the intrinsic KKT certificate (rejected off the optimum).

4 Results

Each use case is pinned to a single channel/source. **G1** validates the bound against Monte-Carlo at small n ; **G2** is the prior gap (optimal achievable prior vs. optimal memoryless vs. the two marginal-

memoryless priors); **G3** compares exact random coding to a closed-form surrogate; **G4** contrasts the error/distortion spectra of the converse- and achievability-optimal priors; and **G5** replaces the (KKT-certified) optimal achievable prior by its i.i.d. *product* (marginalized) version and shows the marginalization cost. The achievability-optimal prior is the KKT-certified march; the bound is the exact random-coding kernel. G1 stays at $n = 8$; G2–G5 are shown at $n = 8$ and $n = 20$. (The converse prior is omitted from G5: its single-threshold optimum is non-unique where the converse is slack, so reusing it for the whole-spectrum integral is ill-posed.)

4.1 Channel coding — $Z(0.1)$

The non-product prior gain peaks at low rate and grows with n ($\approx 2.0\%$ at $n = 8$, $\approx 3.3\%$ at $n = 20$). The converse-optimal prior, reused for achievability, is $5\times$ worse at $n = 8$ and $15\times$ at $n = 20$.




Figure 3: G1 ($n = 8$): random-coding expectation vs. 60 Monte-Carlo codebooks.






Figure 4: G2 — the prior gap (exact kernel). Right panels: gain of the optimal prior over the best memoryless prior.




(a) $n = 8$




(b) $n = 20$

Figure 5: G3 — exact random coding vs. the union bound.





(a) $n = 8$





(b) $n = 20$

Figure 6: G4 — error spectrum of the achievability- vs. converse-optimal prior.



(a) $n = 8$



(b) $n = 20$

Figure 7: G5 — the optimal achievable prior vs. its i.i.d. *product* (marginalized) version (left), and the marginalization cost (right): $\approx 3.5\%$ at low rate, $\rightarrow 0$ at high rate. The optimum is nearly i.i.d.




Figure 8: The complementary view: fix P_e and plot rate vs. blocklength, for two targets $\epsilon = 10^{-2}, 10^{-3}$ (same colour = same ϵ ; solid = converse, dashed = achievable). Both points are a *single solve*: the converse is one LP and the achievable RCU^+ is one convex program (the threshold $w = e^{-R}$ enters as quad-over-linear). They bracket the ϵ -capacity and rise toward C ; the looser $\epsilon = 10^{-2}$ pair is higher and tighter. Capped at $n = 28$ by cvxpy’s compilation ceiling.

4.2 Rate–distortion, average — BMS(0.25) + Hamming

The gain over the best memoryless prior is $\approx 5.4\%$ ($n = 8$) / 5.9% ($n = 20$); the converse and achievability priors are nearly identical ($D : 0.1551 \rightarrow 0.1529$ at $n = 8$).




Figure 9: G1 ($n = 8$): realized distortion vs. the random-coding expectation.




Figure 10: G2 — prior gap (exact best-of- M kernel).




Figure 11: G3 — exact random coding vs. the exponential surrogate.




Figure 12: G4 — distortion spectra (nearly identical priors).




Figure 13: G5 — optimal achievable prior vs. its i.i.d. product: the marginalization cost is $\approx 6.5\%$; the optimum is still essentially i.i.d.






Figure 14: Rate vs. blocklength at fixed distortion $D = 0.1$: the achievable rate converges from above to the *converse*, which for average distortion is the rate-distortion function $R(D) = H(p) - H(D) = 0.342$ (a valid lower bound at every n ; its dispersion term vanishes for average D). The inversion uses the pure-NumPy march (bisection on the rate), which scales to $n = 80$ – past the channel’s single-cvxpy ceiling.

4.3 Rate–distortion, excess — BMS(0.25), $T = 1$, $n = 6$


Excess distortion is the best-of- M of the block indicator, a lifted quantity, so $n = 6$. The gain peaks at $\approx 7\%$; reusing the converse prior costs $2.8\times$ ($P_{\text{exc}} : 4.5\text{e-}2$ vs. $1.6\text{e-}2$).




(a) G1



(b) G2




(c) G3



(d) G4

Figure 15: Excess distortion: bound vs. MC (G1), prior gap (G2), exact vs. bound (G3), and the excess spectra (G4); the converse-marginal prior is the worst here.

G5 RD excess (BMS $p=0.25$, $n=6$, $T=1$): optimal prior vs its product



cost of the i.i.d. product




Figure 16: G5 — optimal achievable prior vs. its i.i.d. product: the marginalization cost is $\approx 8.6\%$ (largest of the three, a probability being more prior-sensitive), but the optimum is still close to i.i.d.

4.4 JSCC — BMS(0.07) over BSC(0.05), $H < C$

With $H = 0.366 < C = 0.714$ bit the source is transmissible. The coded bounds (converse, and the RCU^+ achievable bound at the march optimum) fall with n , while the uncoded symbol-by-symbol scheme degrades; they cross near $n \approx 4$.




Figure 17: JSCC error probability vs. blocklength: meta-converse, achievable (RCU^+ via the simplex march), Monte-Carlo of random codes, and the uncoded baseline.

5 Conclusion

Prior optimization of the achievability bound is one convex program $\max_Q c^\top \Phi(AQ)$; the kernel chooses Φ , and a single first-order simplex march—KKT-certified, exact for any kernel—solves it across channel coding, rate–distortion, and JSCC. With the certified optimum in hand the headline finding is rigorous rather than heuristic: for these minimal settings the memoryless prior is within a few percent of optimal, while the *converse* prior, reused for achievability, can be far from it.